

Quiz 11

Deep Reinforcement Learning Algorithms Comparison

DQN - Deep Q-Network

Characteristics

- Value-based method -> Estimate of the optimal action-value function.
- Off-policy & model-free.
- Neural Network (NN) is used as a function approximator.

Loss Function

$$L = E\left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_k) - Q(s, a; \theta_k)\right)^2\right]$$

Update function

$$\theta_{k+1} = \theta_k + \alpha \left(r + \gamma \max_{a'} Q(s', a'; \theta_k) - Q(s, a; \theta_k) \right) \nabla_{\theta_k} Q(s, a; \theta_k)$$

Cons

- Maximization bias - Selecting the maximum estimated value over and over again causes this bias. This produces low-quality policy & unstable training.

DDQN - Double Deep Q-Network

Characteristics

- Value-based method
- Uses 2 NNs to select & evaluate action.
- We will reduce the *Root Mean Squared* (RMS) Error between the estimated Q & the target Q .

Loss Function

$$L = E\left[\left(r + \gamma \max_{a'} Q(s', a') - Q(s, a)\right)^2\right]$$

Update Function

$$Q_{t+1}^A = (1 - \alpha)Q_t^A(s_t, a_t) + \alpha \left(R_t + \gamma Q_t^B(s_{t+1}, \arg \max_a Q_t^A(s_{t+1}, a)) \right)$$

$$Q_{t+1}^B = (1 - \alpha)Q_t^A(s_t, a_t) + \alpha \left(R_t + \gamma Q_t^B(s_{t+1}, \arg \max_a Q_t^A(s_{t+1}, a)) \right)$$

Algorithm

Algorithm 1 : Double Q-learning (Hasselt et al., 2015)

Initialize primary network Q_θ , target network $Q_{\theta'}$, replay buffer \mathcal{D} , $\tau < 1$

for each iteration do

 for each environment step do

 Observe state s_t and select $a_t \sim \pi(a_t, s_t)$

 Execute a_t and observe next state s_{t+1} and reward $r_t = R(s_t, a_t)$

 Store (s_t, a_t, r_t, s_{t+1}) in replay buffer \mathcal{D}

 for each update step do

 sample $e_t = (s_t, a_t, r_t, s_{t+1}) \sim \mathcal{D}$

 Compute target Q value:

$$Q^*(s_t, a_t) \approx r_t + \gamma Q_{\theta'}(s_{t+1}, \arg \max_{a'} Q_{\theta'}(s_{t+1}, a'))$$

 Perform gradient descent step on $(Q^*(s_t, a_t) - Q_\theta(s_t, a_t))^2$

 Update target network parameters:

$$\theta' \leftarrow \tau * \theta + (1 - \tau) * \theta'$$


Pros

- Avoids maximization bias.
- More stable & reliable than the DQN.

Dueling DQN

Characteristics

- Value-based method
- Has 2 estimators: One for state-value function & other for state-dependent action advantage function.



Aggregate computation

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + (A(s, a; \theta, \alpha) - \frac{1}{A} \sum_{a'} A(s, a'; \theta, \alpha))$$

Pros

- Learns which states are valuable & which are not.
- Fast training.

REINFORCE

Characteristics

- Monte Carlo Policy-gradient method -> Estimate best weight by gradient ascent.
- On-policy
- Updates parameters by stochastic gradient ascent.

Algorithm

REINFORCE, A Monte-Carlo Policy-Gradient Method (episodic)

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$

Repeat forever:

 Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|\cdot, \theta)$

 For each step of the episode $t = 0, \dots, T-1$:

$G \leftarrow$ return from step t

$\theta \leftarrow \theta + \alpha \gamma^t G \nabla_{\theta} \ln \pi(A_t|S_t, \theta)$

Pros

1. Works in environments with discrete or continuous action spaces.

Cons

1. Has high variance.

A2C - Advantage Actor Critic

Characteristics

- Actor-critic (AC) method
 - *Critic*: Updates action-value function.

- *Actor*: Updates policy parameters based on Critic.

• On-policy.

Advantage Function

$$A(s, a) = Q(s, a) - V(s)$$

Pros

1. Advantage function reduces variance.


Cons

1. Not suited for continuous control.

A3C - Asynchronous Advantage Actor Critic

Characteristics

- Uses multiple agents which has its own set of parameters & their own copy of the environment.
- Combines the strength of both *Value-iteration* & *Policy-gradient* methods, and predicts *value function* $V(s)$ & the *optimal policy function* $\pi(s)$.



Algorithm

Algorithm S3 Asynchronous advantage actor-critic - pseudocode for each actor-learner thread.

// Assume global shared parameter vectors θ and θ_v and global shared counter $T = 0$

// Assume thread-specific parameter vectors θ' and θ'_v

Initialize thread step counter $t \leftarrow 1$

repeat

Reset gradients: $d\theta \leftarrow 0$ and $d\theta_v \leftarrow 0$.

Synchronize thread-specific parameters $\theta' = \theta$ and $\theta'_v = \theta_v$

$t_{start} = t$

Get state s_t

repeat

Perform a_t according to policy $\pi(a_t|s_t; \theta')$

Receive reward r_t and new state s_{t+1}

$t \leftarrow t + 1$

$T \leftarrow T + 1$

until terminal s_t **or** $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{ Bootstrap from last state} \end{cases}$

for $i \in \{t-1, \dots, t_{start}\}$ **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt θ' : $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$

Accumulate gradients wrt θ'_v : $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

end for

Perform asynchronous update of θ using $d\theta$ and of θ_v using $d\theta_v$.

until $T > T_{max}$

Pros

1. Copies of agent in environment de-correlates the data.
2. No Experience Replay is needed.

Cons

1. Not suited for continuous control.
2. Data inefficient.

TRPO - Trust Region Policy Optimization**Characteristics**

- Trust Region method
- On-policy [Source] & model-free.
- Adds KL (Kullback-Leibler) divergence for optimization.

Objective Function

$$J(\theta) = E_{s \sim \rho} E_{a \sim \pi_{\theta}^{\text{old}}} \left[\frac{\pi_{\theta}(a|s)}{\pi_{\theta}^{\text{old}}(a|s)} \hat{A}_{\theta^{\text{old}}}(s, a) \right] \text{ or } J^{\text{TRPO}}(\theta) = E[r(\theta) \hat{A}_{\theta^{\text{old}}}(s, a)]$$

Where, $r(\theta)$ is the probability ratio between old & new policies.

Algorithm**Algorithm 1** Trust Region Policy Optimization

- 1: Input: initial policy parameters θ_0 , initial value function parameters ϕ_0
- 2: Hyperparameters: KL-divergence limit δ , backtracking coefficient α , maximum number of backtracking steps K
- 3: **for** $k = 0, 1, 2, \dots$ **do**
- 4: Collect set of trajectories $\mathcal{D}_k = \{\tau_i\}$ by running policy $\pi_k = \pi(\theta_k)$ in the environment.
- 5: Compute rewards-to-go \hat{R}_t .
- 6: Compute advantage estimates, \hat{A}_t (using any method of advantage estimation) based on the current value function V_{ϕ_k} .
- 7: Estimate policy gradient as

$$\hat{g}_k = \frac{1}{|\mathcal{D}_k|} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t|s_t)|_{\theta_k} \hat{A}_t.$$

- 8: Use the conjugate gradient algorithm to compute

$$\hat{x}_k \approx \hat{H}_k^{-1} \hat{g}_k,$$

where \hat{H}_k is the Hessian of the sample average KL-divergence.

- 9: Update the policy by backtracking line search with

$$\theta_{k+1} = \theta_k + \alpha^j \sqrt{\frac{2\delta}{\hat{x}_k^T \hat{H}_k \hat{x}_k}} \hat{x}_k,$$

where $j \in \{0, 1, 2, \dots, K\}$ is the smallest value which improves the sample loss and satisfies the sample KL-divergence constraint.

- 10: Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left(V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

- 11: **end for**

Pros

- Suitable for environments with both continuous & discrete action space.

Cons

- Data inefficient

PPO - Proximal Policy Optimization

Characteristics

- Trust Region method
- On-policy [Source 1 & Source 2] & model-free.
- Learns policy & value function at the same time.
- Improves/simplifies on TRPO by adding a *clipped surrogate objective*.
- Has entropy component which is the measure of uncertainty in the policy (i.e.) lower the entropy, the policy is more confident in choosing an action.
- Prefers exploration & avoids bad local optimum by rewarding for choosing actions with high entropy.

Objective Function

$$J^{\text{CLIP}}(\theta) = \hat{E}_t \left[\min \left(r(\theta) \hat{A}_{\theta_{\text{old}}}(s, a), \text{clip}(r(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_{\theta_{\text{old}}}(s, a) \right) \right]$$

The clip component clips the ratio within $[1 - \epsilon, 1 + \epsilon]$.

Algorithm

Algorithm 5 PPO with Clipped Objective

Input: initial policy parameters θ_0 , clipping threshold ϵ

for $k = 0, 1, 2, \dots$ **do**

Collect set of partial trajectories \mathcal{D}_k on policy $\pi_k = \pi(\theta_k)$

Estimate advantages $\hat{A}_t^{\pi_k}$ using any advantage estimation algorithm

Compute policy update

$$\theta_{k+1} = \arg \max_{\theta} \mathcal{L}_{\theta_k}^{\text{CLIP}}(\theta)$$

by taking K steps of minibatch SGD (via Adam), where

$$\mathcal{L}_{\theta_k}^{\text{CLIP}}(\theta) = \mathbb{E}_{\tau \sim \pi_k} \left[\sum_{t=0}^T \left[\min(r_t(\theta) \hat{A}_t^{\pi_k}, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t^{\pi_k}) \right] \right]$$

end for

Pros

1. Simpler compared to TRPO, but same performance.
2. Works in both discrete & continuous environments.
3. Faster & stable training.
4. Entropy regularization

Cons

1. Data inefficient

SAC - Soft Actor Critic

Characteristics

- Off-policy
- Avoids convergence to bad local optimum by rewarding actions with high entropy.
- Learns value function & the policy at the same time.

Objective function

$$J(\pi) = \sum_{t=0}^T E_{(s_t, a_t) \sim \rho_{\pi}} \left[(s_t, a_t) + \alpha H(\pi(\cdot | s_t)) \right]$$

Algorithm

Algorithm 1 Soft Actor-Critic

- 1: Input: initial policy parameters θ , Q-function parameters ϕ_1, ϕ_2 , empty replay buffer \mathcal{D}
- 2: Set target parameters equal to main parameters $\phi_{\text{targ},1} \leftarrow \phi_1, \phi_{\text{targ},2} \leftarrow \phi_2$
- 3: **repeat**
- 4: Observe state s and select action $a \sim \pi_\theta(\cdot|s)$
- 5: Execute a in the environment
- 6: Observe next state s' , reward r , and done signal d to indicate whether s' is terminal
- 7: Store (s, a, r, s', d) in replay buffer \mathcal{D}
- 8: If s' is terminal, reset environment state.
- 9: **if** it's time to update **then**
- 10: **for** j in range(however many updates) **do**
- 11: Randomly sample a batch of transitions, $B = \{(s, a, r, s', d)\}$ from \mathcal{D}
- 12: Compute targets for the Q functions:

$$y(r, s', d) = r + \gamma(1 - d) \left(\min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

- 13: Update Q-functions by one step of gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$

- 14: Update policy by one step of gradient ascent using

$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} \left(\min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_\theta(s)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s)|s) \right),$$

where $\tilde{a}_\theta(s)$ is a sample from $\pi_\theta(\cdot|s)$ which is differentiable wrt θ via the reparametrization trick.

- 15: Update target networks with

$$\phi_{\text{targ},i} \leftarrow \rho \phi_{\text{targ},i} + (1 - \rho) \phi_i \quad \text{for } i = 1, 2$$

- 16: **end for**
 - 17: **end if**
 - 18: **until** convergence
-

Pros

1. Sample efficient - Useful when environment is expensive to sample from.
2. Entropy maximization - Data efficient when compared to ER of PPO & balances exploitation-exploration.

Cons

1. Suitable only for environment with continuous action space.
2. Unstable when compared to PPO.